


# Ray tracing



ING 

# Ray tracing pseudo-algorithm

```
for pixel in pixels {  
    pixel.setBlack()  
    minDistance = Infinity  
    indexMin = -1  
    for sphere in spheres {  
        if (sphere.isBelow(pixel) && sphere.distance(pixel) < minDistance) {  
            minDistance = sphere.distance()  
            indexMin = sphereIndex  
        }  
    }  
    if (indexMin != -1) {  
        pixel.setHue(spheres[indexMin].getHue())  
        pixel.setBrightness(spheres[indexMin].getBrightness())  
    }  
}
```

# Time measures definitions

- CPU
  - Init, memory allocation (spheres, image)
  - Compute, fillImageGL
- GPU
  - Init, host memory allocation + host->device copy
  - Compute, kernel fillImageGL
  - GetResult, copy image device->host

# Time measures definitions



- CPU
  - Init, memory allocation (spheres, image)
  - Compute, fillImageGL
- GPU
  - Init, host memory allocation + host->device copy
  - Compute, kernel fillImageGL
  - GetResult, copy image device->host

# Time complexity

- Program version used
  - CPU
  - Intel compiler
  - OMP 1 thread => 1 CPU core
- Init complexity
- Compute complexity


# Ray tracing reminder

- Two size parameters
  - Number of spheres
  - Number of pixels
- How the two parameters influence time complexity ?



# Init time complexity discussion

- Why init time complexity seems to be constant,  $O(1)$ , versus number of sphere ?




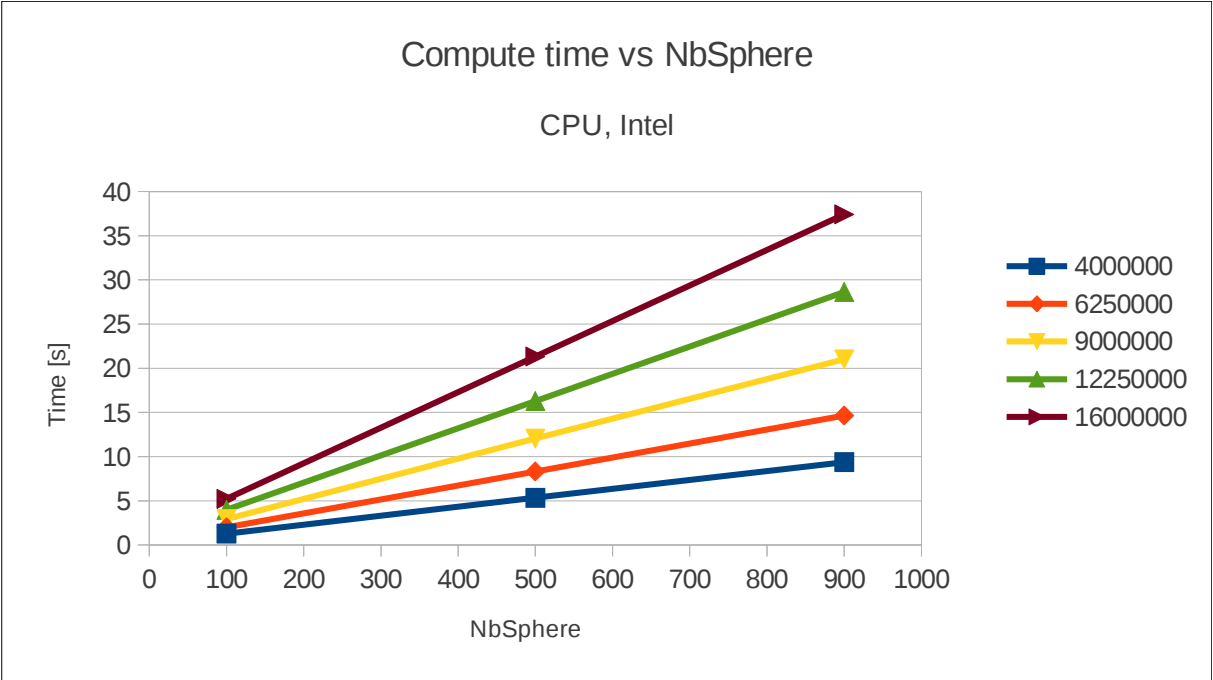


# Init time complexity discussion (2)

- Why init time complexity seems to be constant,  $O(1)$ , versus number of sphere ?
  - Same measure for two different allocations:
    - Image (pixels)
    - Spheres
  - The memory allocation of spheres is not constant, but less significant than the pixel memory allocation, because there is a lot more of pixels than spheres.

# Init time complexity conclusions

- One measure for two different things could lead to false interpretation.
- Init time versus number of pixel follows a linear complexity  $O(n)$ .
- Init time versus number of spheres should follow a linear complexity.
  - Measures with greater numbers of spheres are required to determine this hypothesis.



# Compute time complexity conclusions


- Both parameters induce a linear complexity.
- Slope of linear law grows as second “size” parameter grows.
- Could be interesting to plot the evolution of slope versus second parameter.

# CPU – Compiler comparison

- Compilers
  - Intel
  - MinGW
  - Visual
- Benchmark parameters
  - NbPixel =  $2000 \times 2000 = 4000000$
  - NbSphere = {100, 500, 900}
  - NbThread = {1, 2, 3, 4, 5}


### Speedup OMP Intel

NbPixel = 4000000




### Speedup OMP MinGW

NbPixel = 4000000




### Speedup OMP Visual

NbPixel = 4000000



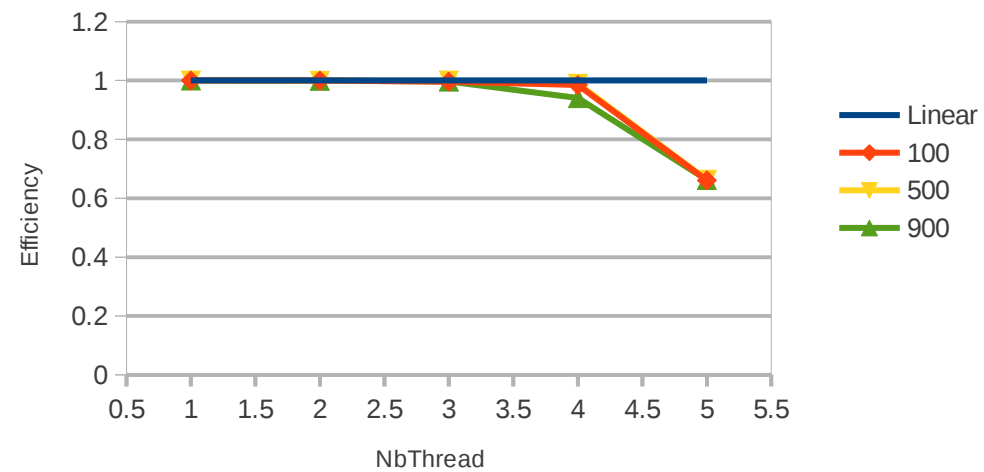
### Efficiency OMP Intel

NbPixel = 4000000




### Efficiency OMP MinGW

NbPixel = 4000000



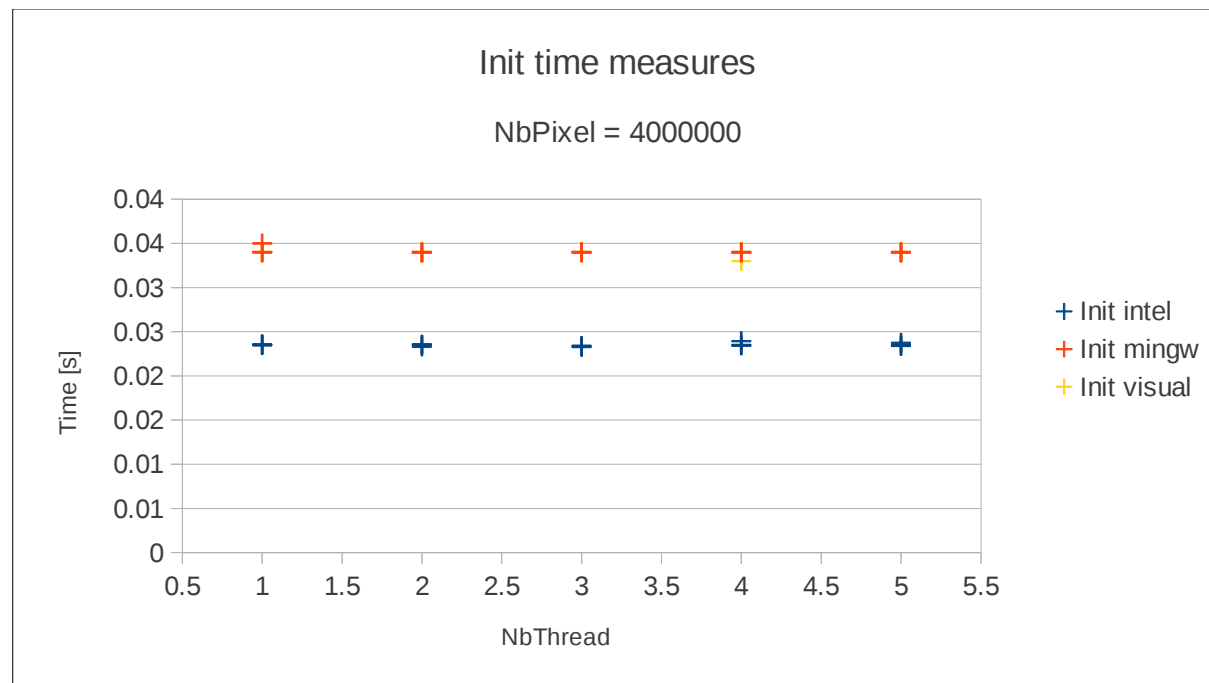
### Efficiency OMP Visual

NbPixel = 4000000



# Init time comparison

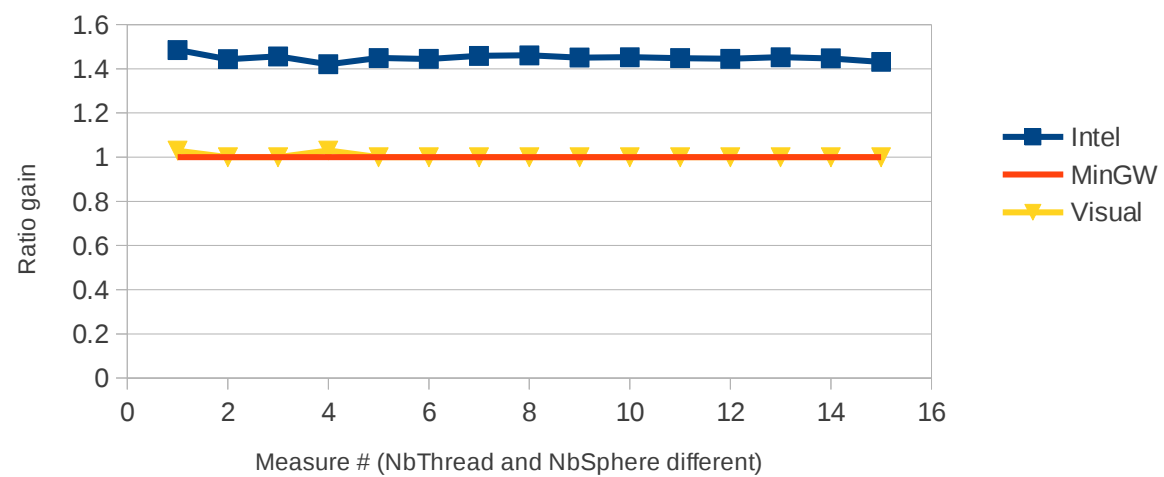
- $3 \text{ (NbSphere)} * 5 \text{ (NbThread)} = 15 \text{ measures}$
- 15 measures for each compiler for a given number of pixel





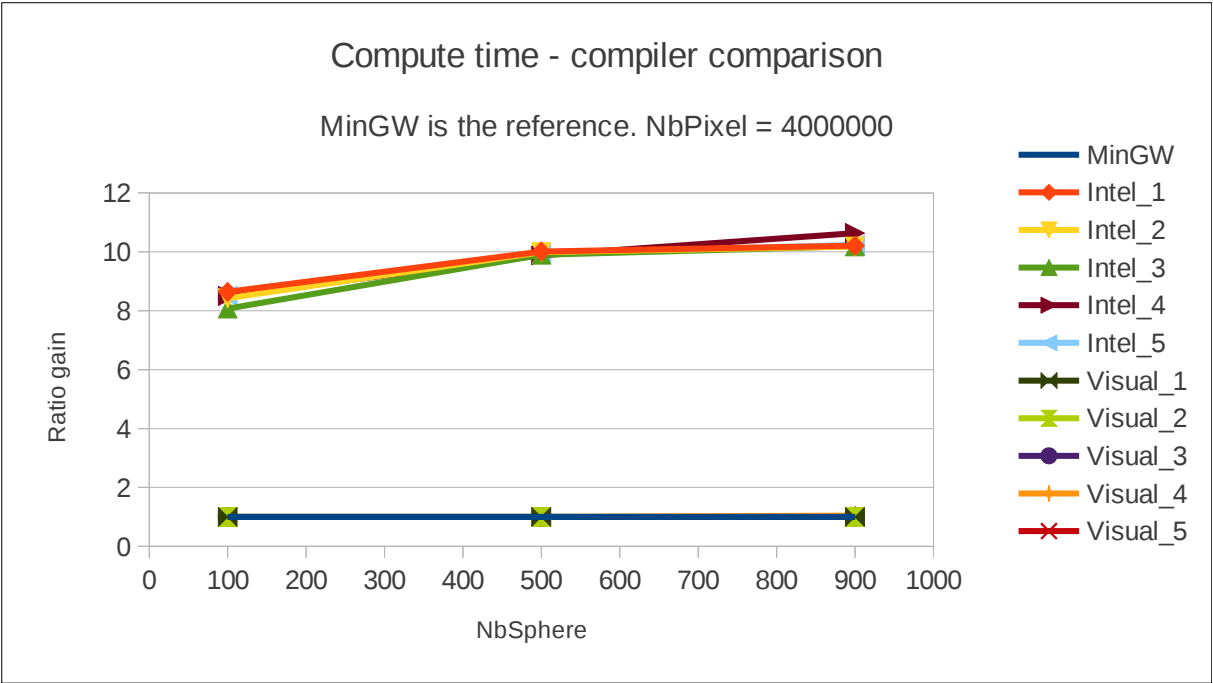

### Init time measure comparison

MinGW is the reference. NbPixel = 4000000



# Conclusion of init time comparison

- Intel compiled code is at least 1.4 time faster for memory allocation (pixel mostly) than MinGW or Visual.
- Visual and MinGW compilers have equivalent performance (or are as bad) for memory allocation.
- Init time is not improved by number of thread because it is a sequential code.



# Conclusion of compiler comparison for compute time

- Performance gain provided by Intel compiler is independent from number of thread.
- Intel performance gain versus number of sphere follows a logarithmic law.
- Visual and MinGW compilers are as bad as one another for computation.
- Intel compiler is up to ten times faster than Visual or MinGW for computation part.

# GPU

- Benchmarks parameters
  - NbPixel = [1'000'000, 100'000'000]
  - NbSphere = 500
  - MemType = {Global, Shared, Constant}
  - dg (dimension of grid) and db (dimension of block) defined by range later.

# GPU, wrap a present

- Could a wrap be made of threads from different blocks ?
  - If yes, we could assume that in term of performance  $16 \times 32$  and  $512 \times 1$  (dg x db) are closed (for none-shared versions).
  - If no, performance of  $512 \times 1$  should be horrible compared to  $16 \times 32$ .

# GPU, push the limits

- Could we measure the speedup of streaming multiprocessors ?
  - By design, a block is “assigned” to a SM =>  $dg = [1, 16]$ .
  - $dg$  becomes the correspondence of `nbThread` in OMP.
  - **Total number of thread is kept constant**
    - Multiple wraps per SM when  $dg < 16$ .
    - Could we obtain “good” performance by using  $dg < 16$  and  $db > 32$  ?
    - What loss to expect when  $dg > 16$  and  $db < 32$  ?
  - Total number of thread is variant
    - Dimension of block is a constant,  $db = 32$ .
    - Total number of thread depends of  $dg$ .
    - Would better correspond to the speedup we know with CPU cores.

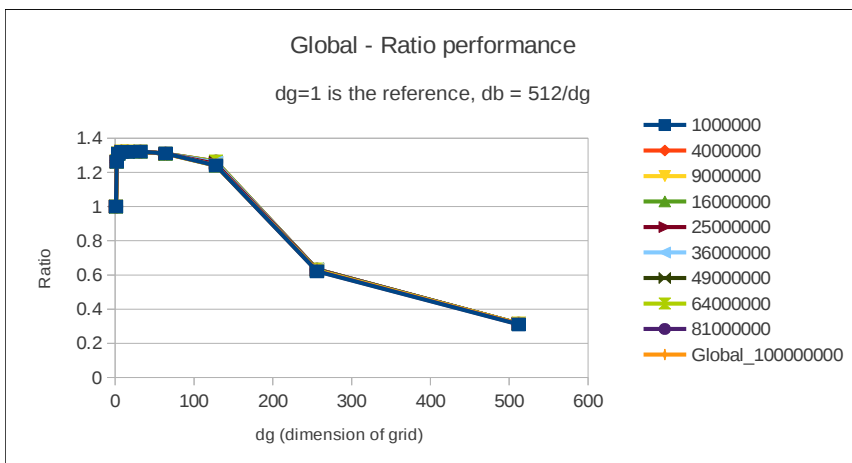
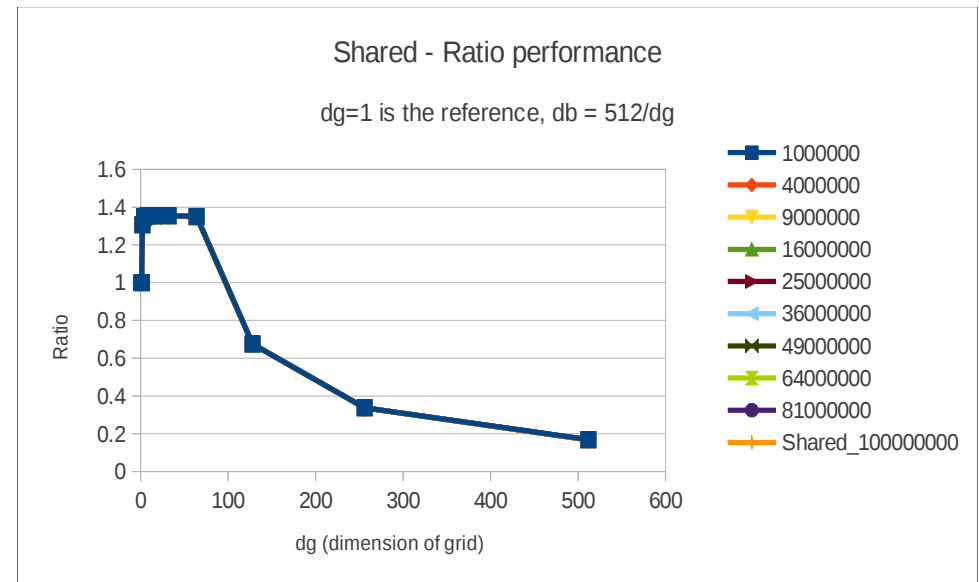
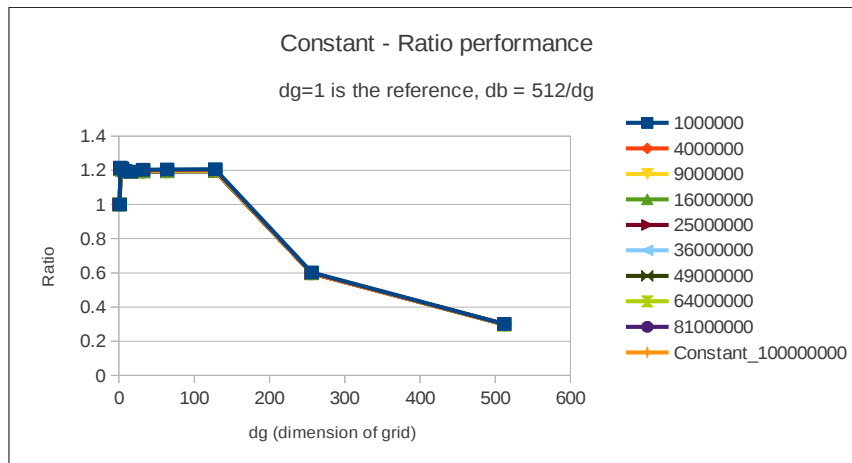
# Parameters dg,db

- $dg = [1, 512]$
- $db = [1, 512]$
- $dg * db = 512$
- $dg,db = \{1,512; 2,256; 4,128; 8,64; 16,32; 32,16; 64,8; 128,4; 256,2; 512,1\}$



# Basic observations

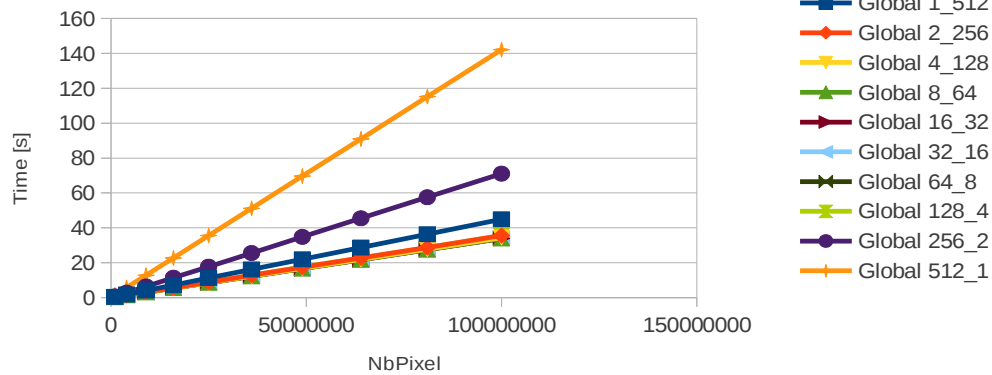
- Performance gain/loss by changing  $dg, db$  is independent of number of pixel.
- Performance gain/loss is dependent of memory type.



# Basic observations


## Global memory - Compute time versus NbPixel

nbSphere = 500



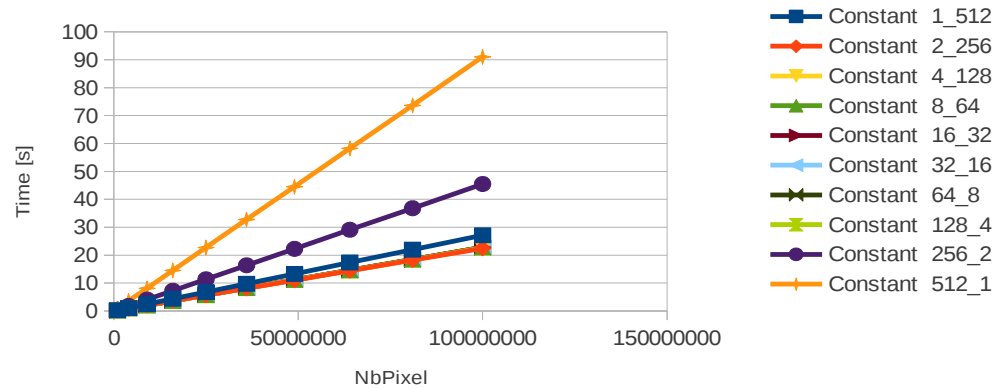
### Shared memory - Compute time versus NbPixel

nbSphere = 500




### Constant memory - Compute time versus NbPixel

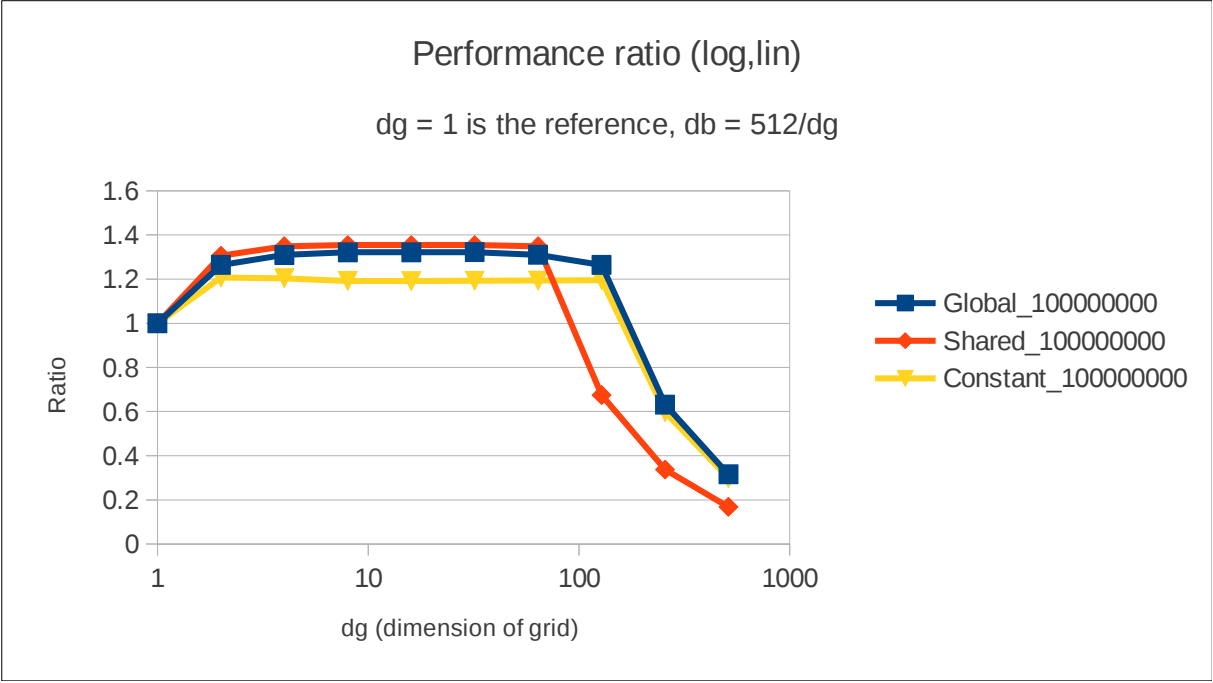
nbSphere = 500



### Performance ratio

$dg = 1$  is the reference,  $db = 512/dg$









# Conclusion GPU dg,db

- First of all let's define "good" performance as a performance gain similar of 16x32 (+- 0.1).
- $2 \leq dg \leq 128$  ( $4 \leq db \leq 256$ )
  - -> performance are "good" for constant and global memory.
- $2 \leq dg \leq 64$  ( $8 \leq db \leq 256$ )
  - -> performance are "good" for shared memory.
- With  $dg < 16$  (and thus  $db > 32$ ), "good" performance are achievable.
  - In other words: We can achieve good performance without using all SM.
- Loss when  $16 < dg < 64$  (and thus  $db < 32$ !) is mitigate.
  - In other words: We can achieve good performance with wrap  $< 32$ .
- Value of  $dg > 64$  ( $db < 8$ ) induces great loss.



# Rule them all

- $db = [2^3, 2^{15}]$
- $dg = [2^3, 2^{15}]$
- $NbSphere = 500$
- $NbPixel = 25'000'000$

# Rule them all Shared





# Rule them all Global





# Rule them all Constant



# dg vs db

- Specific for each algorithm
  - And memory implementation
- Apply rules to “stuff it”
  - dg should be a multiple of SM
  - db should be a multiple of wrap size
- Compute dg,db by benchmarking or theory
- Greater dg is for sure useless without a sufficient db

# dg vs db

- Specific for each algorithm
  - And memory implementation
- Apply rules to “stuff it”
  - dg should be a multiple of SM
  - db should be a multiple of wrap size
- Compute dg,db by benchmarking or theory
- Greater dg is for sure useless without a sufficient db

# Hey listener are you attentive ?


- Why all this deep blue in preceding plots ?
- Is there upper limits for  $db, dg$  ?

# Hey listener are you attentive ?

- Why all this deep blue in preceding plots ?
- Is there upper limits for db,dg ?
  - Yes they are upper limits:
    - Maximum x- or y-dimension of a block 1024
    - Maximum number of threads per block 1024
- Kernel call with bad dimension without check for errors → data in plot

# Memory comparison


- Performance gain of constant over global is less affected by  $dg, db$  values than shared over global.
- Shared is better for normal values of  $dg, db$ 
  - Normal = multiplier of SM/wrap, 16x32 in this case




# Global, shared or constant

- Each implementation is affected differently by the dg,db parameters
- Thus, each implementation has a different dg,db optimized value.


# Shared



# Global




# Constant



## Memory comparison

Global memory is the reference,  $dg = 2^8 = 256$





# Final conclusion

- Is the “stuff it” a best practice ?
- Concurrent kernel with Fermi architecture.
- The search of perfect dg,db.
- The others dimensions of dg,db.
- Kernel side measures (memory copies)
- The curious spot in dg,db plots.